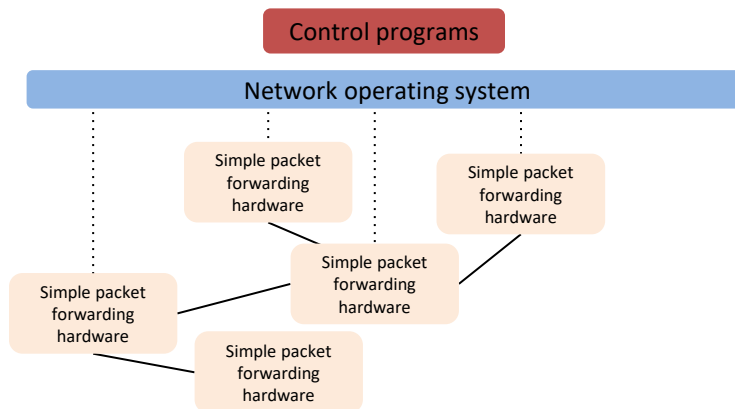


Software Defined Networking (SDN)

- Separation of control functions from forwarding functions
- Intelligence is concentrated in the control plane



1

Hands-On Experimentation With SDN

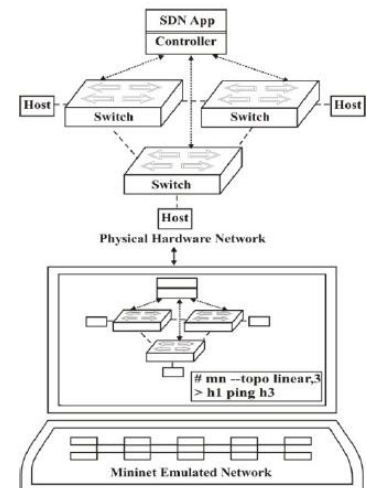
- Need a capability (simulator or emulator) to create real-world scenarios
- Mininet is an emulator that provides such capabilities
- Mininet uses the container-based abstractions
 - Namespaces, cgroups
 - Can create arbitrary topologies
 - Can emulate hosts and hosts can execute user-supplied logic
 - Supports open virtual switch
 - Implements OpenFlow controller
- On your laptop VM, do:
 - Clone the Mininet github repo and invoke the install.sh command with -a flag as discussed in the next few slides.
 - There is also a special Mininet VM available

2

Mininet: Basic Usage, CLI, API

Mininet: network emulator which creates realistic virtual network

- Runs real kernel, switch, and application code on a single machine
- Provides both Command Line Interface (CLI) and Application Programming Interface (API)
- Reasonably accurate, easy to download, and fast/interactive usage
- Abstraction
 - Host: emulated as an OS level process
 - Switch: emulated by using software-based switch
- **But, slower than hardware experiments: may not fit possible inaccuracy from multiplexing**



3

Mininet: Basic Usage, CLI, API (cont.)

Install Mininet (multiple different approaches; I used Git and install)

- **Mininet VM installation**
 - Download the [Mininet VM image](#)
- **Native installation from source (Preferred Approach)**
 - `git clone git://github.com/mininet/mininet`
 - `$ mininet/util/install.sh -a`
 - You may need to do “`sudo apt-get install net-tools`” for missing `ifconfig`
- **Installation from packages**
 - `$ sudo apt-get install mininet`
- **Upgrading an existing Mininet installation**
 - `$ cd mininet && git fetch && git checkout master number or a specific version like 2.2.1`
 - `$ git pull && sudo make install`

Mininet uses Linux network namespaces and some additional capabilities.

4

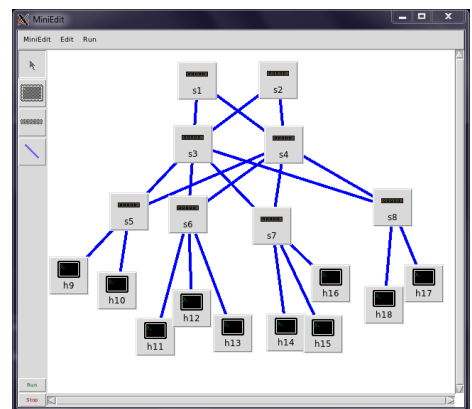
Neat Features of Mininet

- Ability to create arbitrarily large and complex networking topologies
- Ability to execute application logic in emulated hosts
- Ability to control network traffic characteristics
- Ability to integrate with Wireshark to observe traffic

5

Mininet: Basic Usage, CLI, API

- **Mininet GUI**
 - Mininet provides a tool (“MiniEdit”) to help creating real networks
 - Run the file mininet/examples/miniedit.py
- **Mininet CLI and features**
 - `$sudo mn --topo tree,depth=3,fanout=3 --link=tc,bw=10`
 - `mininet> xterm h1 h2`
- **Mininet’s Python API**
 - `net = Mininet()` # net is a Mininet() object
 - `h1 = net.addHost('h1')` # h1 is a Host() object
 - `h2 = net.addHost('h2')` # h2 is a Host()
 - `s1 = net.addSwitch('s1')` # s1 is a Switch() object



6

Simple Command to Try

```
gokhale@gokhale-ubuntu: ~
gokhale@gokhale-ubuntu:~$ sudo mn
*** No default OpenFlow controller found for default switch!
*** Falling back to OVS Bridge
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> 
```

Simplest case has two hosts and one switch. Mininet will create the topology including assigning IP addresses, etc.

CLI where we can try out different commands

7

Simple Command to Try (cont.)

```
mininet>
mininet> nodes
available nodes are:
h1 h2 s1
mininet>
mininet>
mininet>
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0
mininet>
mininet>
mininet>
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=6772>
<Host h2: h2-eth0:10.0.0.2 pid=6775>
<OVSBridge s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=6781>
mininet>
mininet>
mininet> 
```

Commands to try on the CLI:
Use help to find more commands.

8

Mininet Help on the CLI

```
mininet>
mininet> help

Documented commands (type help <topic>):
=====
EOF      gterm  iperfudp  nodes      pingpair    py      switch
dpctl    help   link      noecho     pingpairfull  quit    time
dump     intf  links     pingall    ports       sh      x
exit     iperf  net       pingallfull  px          source  xterm

You may also send a command to a node using:
  <node> command {args}
For example:
  mininet> h1 ifconfig

The interpreter automatically substitutes IP addresses
for node names when a node is the first arg, so commands
like
  mininet> h2 ping h3
should work.

Some character-oriented interactive commands require
noecho:
  mininet> noecho h2 vi foo.py
```

9

Command Line Parameters to “sudo mn”

```
gokhale@gokhale-ubuntu:~$ sudo mn -h
Usage: mn [options]
(type mn -h for details)

The mn utility creates Mininet network from the command line. It can create
parametrized topologies, invoke the Mininet CLI, and run tests.

Options:
  -h, --help                show this help message and exit
  --switch=SWITCH            default|ivs|lxb|ovs|ovsbr|ovsk[user[,param=value...]]
                             ovs=OVSSwitch default=OVSSwitch ovsk=OVSSwitch
                             lxb=Linux8Ridge user=UserSwitch ivs=IVSSwitch
                             ovsbr=OVSBridge
  --host=HOST                cfs|proc|rt[,param=value...]
                             rt=CPULimitedHost('sched': 'rt') proc=Host
                             cfs=CPULimitedHost('sched': 'cfs')
  --controller=CONTROLLER   default|none|nox|ovsc|ref|remote|ryu[,param=value...]
                             ovsc=OVSController none=NullController
                             remote=RemoteController default=DefaultController
                             nox=NOX ryu=Ryu ref=Controller
  --link=LINK                default|ovs|tc[,param=value...] default=Link
                             ovs=OVSLink tc=TCLink
  --topo=TOPO                linear|minimal|reversed|single|torus|tree[,param=value
                             ...] linear=LinearTopo torus=TorusTopo tree=TreeTopo
                             single=SingleSwitchTopo
                             reversed=SingleSwitchReversedTopo minimal=MinimalTopo
  -C, --clean                clean and exit
  --custom=CUSTOM            read custom classes or params from .py file(s)
  --test=TEST                cli|build|pingall|pingpair|iperf|all|iperfudp|none
  -x, --xterms               spawn xterms for each node
  -i IPBASE, --ipbase=IPBASE base IP address for hosts
  --mac                      automatically set host MACs
  --arp                      set all-pairs ARP entries
  -v VERBOSITY, --verbosity=VERBOSITY
```

10

Executing Commands on Hosts, Switches, and Controllers

```

mininet> h1 ifconfig -a
h1-eth0 Link encap:Ethernet HWaddr aa:dc:88:0vs-system
        inet addr:10.0.0.1 Bcast:10.255.255.
        inet6 addr: fe80::a8dc:8ff:feac:93b3/
        UP BROADCAST RUNNING MULTICAST MTU:1
        RX packets:30 errors:0 dropped:0 over
        TX packets:8 errors:0 dropped:0 overr
        collisions:0 txqueuelen:1000
        RX bytes:3225 (3.2 KB) TX bytes:648
Link encap:Local Loopback
        inet addr:127.0.0.1 Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
        UP LOOPBACK RUNNING MTU:65536 Metri
        RX packets:0 errors:0 dropped:0 overr
        TX packets:0 errors:0 dropped:0 overr
        collisions:0 txqueuelen:1
        RX bytes:0 (0.0 B) TX bytes:0 (0.0 B) s1

mininet> h2 ifconfig -a
h2-eth0 Link encap:Ethernet HWaddr 66:8d:a7:
        inet addr:10.0.0.2 Bcast:10.255.255.
        inet6 addr: fe80::668d:a7ff:fe2c:6948
        UP BROADCAST RUNNING MULTICAST MTU:1
        RX packets:27 errors:0 dropped:0 overr
        TX packets:8 errors:0 dropped:0 overr
        collisions:0 txqueuelen:1000
        RX bytes:3135 (3.1 KB) TX bytes:648
Link encap:Local Loopback
        inet addr:127.0.0.1 Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
        UP LOOPBACK RUNNING MTU:65536 Metri
        RX packets:0 errors:0 dropped:0 overr
        TX packets:0 errors:0 dropped:0 overr
        collisions:0 txqueuelen:1
        RX bytes:0 (0.0 B) TX bytes:0 (0.0 B) s1-eth2

mininet> h1 ps
PID TTY TIME CMD
8426 pts/17 00:00:00 bash
8781 pts/17 00:00:00 ps
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.259 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.135 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.124 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.104 ms

```

Output of s1 ifconfig -a

11

Executing Commands on Hosts, Switches, and Controllers (cont.)

- To execute a specific command inside a given entity (e.g., a host or switch), specify the entity name first, followed by the command you wish to execute.
- Example here shows running the ifconfig command, ps command, pinging other host in the network, etc.

```

mininet> h1 ifconfig -a
h1-eth0 Link encap:Ethernet HWaddr da:c0:8b:04:2b:7e
        inet addr:10.0.0.1 Bcast:10.255.255.255 Mask:255.0.0.0
        inet6 addr: fe80::d8c0:8bff:fe04:2b7e/64 Scope:Link
        UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
        RX packets:30 errors:0 dropped:0 overruns:0 frame:0
        TX packets:8 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:1000
        RX bytes:3439 (3.4 KB) TX bytes:648 (648.0 B)

Link encap:Local Loopback
        inet addr:127.0.0.1 Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
        UP LOOPBACK RUNNING MTU:65536 Metric:1
        RX packets:0 errors:0 dropped:0 overruns:0 frame:0
        TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:1
        RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

mininet> h1 ps
PID TTY TIME CMD
8426 pts/17 00:00:00 bash
8781 pts/17 00:00:00 ps
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.259 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.135 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.124 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.104 ms



```

12

Mininet: Slightly Advanced Usage, CLI, API

Example 1: create a network and run a simple performance test

- **Minimal topology:** `$sudo mn --topo single,3`
 - Minimal is very simple topology that contains one OpenFlow switch and three hosts.
- **Reversed topology:** `$sudo mn --topo reversed,4`
 - It is similar to single topology but connection order is reversed.
- **Linear topology:** `$sudo mn --topo linear,3`
 - Linear topology contains k switches and k hosts.
 - It also creates a link between each switch and each host and among the switches.




13

Mininet: Slightly Advanced Usage, CLI, API (cont.)

Example 2: create a network and run a simple performance test

- **Tree topology:** `$sudo mn --topo tree,3`
 - Tree topology contains k levels and two hosts are attached to per switch.



14

Namespaces in Mininet

- Note that Mininet uses Linux namespaces for isolating the hosts, switches, and controllers from each other.
- Note that it uses only the network namespace and nothing else.
 - The process hierarchy and file system is the same across all hosts and switches and controllers.
 - One could use additional isolation dimensions if needed using the specific option provided by Mininet.
 - Option to use: `--innamespace`

15

Cleaning Up

- Type `exit` on the mininet CLI to tear down the network and cleanup
- On an exception or if mininet crashes:
 - `sudo mn -c`

16

Mininet: More Complex Example (Example 3)

- `sudo mn --link tc,bw=10,delay=10ms`
- Two hosts and one switch scenario
- Each link has 10 msec delay and a 10 Mbps capacity
- Host h1 ping to h2 will traverse a path from h1 to s1 to h2 and back to s1 to h1 for a total of 40 ms round-trip
- Notice how the first ping took more time because of route setup time

```
gokhale@gokhale-ubuntu:~$ sudo mn --link tc,bw=10,delay=10ms
*** No default OpenFlow controller found for default switch!
*** Falling back to OVS Bridge
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(10.00Mbit 10ms delay) *** Error: RTNETLINK answers: No such file or directory
(10.00Mbit 10ms delay) *** Error: RTNETLINK answers: No such file or directory
(h1, s1) (10.00Mbit 10ms delay) *** Error: RTNETLINK answers: No such file or di
rectory
(h2, s1) (10.00Mbit 10ms delay) *** Error: RTNETLINK answers: No such file or directory
*** Configuring hosts
h1 h2
*** Starting controller
*** Starting 1 switches
s1 ... (10.00Mbit 10ms delay) (10.00Mbit 10ms delay)
*** Starting CLI:
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp seq=1 ttl=64 time=99.2 ms
64 bytes from 10.0.0.2: icmp seq=2 ttl=64 time=41.4 ms
64 bytes from 10.0.0.2: icmp seq=3 ttl=64 time=42.4 ms
64 bytes from 10.0.0.2: icmp seq=4 ttl=64 time=41.8 ms
64 bytes from 10.0.0.2: icmp seq=5 ttl=64 time=41.9 ms
^C
--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4004ms
rtt min/avg/max/mdev = 41.429/53.380/99.234/22.929 ms
mininet>
```

17

Mininet: Programming API: Part I

Example 4: create a customized network topology

- Using Mininet CLI: `$sudo mn --topo tree,depth=3,fanout=3 --link=tc,bw=10`
- Needs just writing a few lines of Python

```
#!/usr/bin/python
from mininet.topo import Topo
from mininet.net import Mininet
from mininet.node import CPULimitedHost
from mininet.link import TCLink
from mininet.util import dumpNodeConnections
from mininet.log import setLogLevel
```

```
class SingleSwitchTopo(Topo):
```

```
    "Single switch connected to n hosts."
```

```
    def __init__(self, n=2, **opts):
```

```
        Topo.__init__(self, **opts)
```

```
        switch = self.addSwitch('s1')
```

```
        for h in range(n):
```

```
            #Each host gets 50%/n of system CPU
```

```
            host = self.addHost('h%s' % (h + 1), cpu=.5/n)
```

```
            #10 Mbps, 5ms delay, 0% Loss, 1000 packet queue
```

```
            self.addLink(host, switch, bw=10, delay='5ms',
```

```
            loss=0, max_queue_size=1000, use_htb=True)
```

```
def perfTest():
```

```
    "Create network and run simple performance test"
```

```
    topo = SingleSwitchTopo(n=4)
```

```
    net = Mininet(topo=topo, host=CPULimitedHost, link=TCLink)
```

```
    net.start()
```

```
    print "Dumping host connections"
```

```
    dumpNodeConnections(net.hosts)
```

```
    print "Testing network connectivity"
```

```
    net.pingAll()
```

```
    print "Testing bandwidth between h1 and h4"
```

```
    h1, h4 = net.get('h1', 'h4')
```

```
    net.iperf(h1, h4)
```

```
    net.stop()
```

```
if __name__ == '__main__':
```

```
    setLogLevel('info')
```

```
    perfTest()
```

18

Mininet: Programming API: Part II

Example 5: create a simple network and use a POX controller

Note that POX, Ryu etc are old SDN controllers; Need to switch to P4, etc

```
#!/usr/bin/python

from mininet.net import Mininet
from mininet.node import Controller
from mininet.topo import SingleSwitchTopo
from mininet.log import setLogLevel

import os

class POXBridge( Controller ):
    "Custom Controller class to invoke POX forwarding.I2_learning"
    def start( self ):
        "Start POX learning switch"
        self.pox = '%s/pox/pox.py' % os.environ[ 'HOME' ]
        self.cmd( self.pox, 'forwarding.I2_learning &' )
    def stop( self ):
        "Stop POX"
        self.cmd( 'kill %' + self.pox )

controllers = { 'poxbridge': POXBridge }

if __name__ == '__main__':
    setLogLevel( 'info' )
    net = Mininet( topo=SingleSwitchTopo( 2 ), controller=POXBridge )
    net.start()
    net.pingAll()
    net.stop()

$ sudo mn --custom poxbridge.py --controller poxbridge --topo tree,2,2 --test pingall -v output
```

19

Mininet: Programming API: Part III

Example 6: dynamically change the network parameters—change link delay

```
#!/usr/bin/python

from mininet.net import Mininet
from mininet.node import Node
from mininet.link import *
from mininet.log import setLogLevel, info
from threading import Timer

def myNet():
    "Create network from scratch using Open vSwitch."
    info( "**** Creating nodes\n" )
    switch0 = Node( 's0', inNamespace=False )
    switch1 = Node( 's1', inNamespace=False )
    h0 = Node( 'h0' )
    h1 = Node( 'h1' )
    info( "**** Creating links\n" )
    linkopts0=dict(bw=100, delay= '1ms', loss=0)
    link1=TCLink( h0, switch0, **linkopts0)
    TCLink( switch0, switch1, **linkopts0)
    link2=TCLink( h1, switch4, **linkopts0)
    info( "**** Configuring hosts\ns" )
    h0.setIP( '192.168.123.1/24' )

h1.setIP( '192.168.123.2/24' )
info( str( h0 ) + '\n' )
info( str( h1 ) + '\n' )

info( "**** Starting network using Open vSwitch\n" )
switch0.cmd( 'ovs-vsctl del-br dp0' )
switch0.cmd( 'ovs-vsctl add-br dp0' )
switch1.cmd( 'ovs-vsctl del-br dp1' )
switch1.cmd( 'ovs-vsctl add-br dp1' )

for intf in switch0.intfs.values():
    print intf
    print switch0.cmd( 'ovs-vsctl add-port dp0 %s' % intf )

for intf in switch1.intfs.values():
    print intf
    print switch1.cmd( 'ovs-vsctl add-port dp1 %s' % intf )
```

20

Mininet: Programming API: Part IV

Example 7: dynamically change the network parameters—change link delay

```
print switch1.cmd('ovs-ofctl add-flow dp1 idle_timeout=0,priority=1,in_port=1,actions=flood')
print switch1.cmd('ovs-ofctl add-flow dp1 idle_timeout=0,priority=1,in_port=1,actions=output:2')
print switch0.cmd('ovs-ofctl add-flow dp0 idle_timeout=0,priority=10,ip,nw_dst=192.168.123.2,nw_tos=0x10,actions=output:2')
print switch0.cmd('ovs-ofctl add-flow dp0 idle_timeout=0,priority=10,ip,nw_dst=192.168.123.2,nw_tos=0x20,actions=output:3')
print switch0.cmd('ovs-ofctl add-flow dp0 idle_timeout=0,priority=10,ip,nw_dst=192.168.123.2,nw_tos=0x30,actions=output:4')

def cDelay1():
    h1.cmdPrint('ethtool -K h1-eth0 gro off')
    h1.cmdPrint('tc qdisc del dev h1-eth0 root')
    h1.cmdPrint('tc qdisc add dev h1-eth0 root handle 10: netem delay 100ms')


def hello(a, b='4'):
    print a
    print b
t=Timer(5.0, cDelay1)
t.start()
```

```
info( "**** Running test\n" )
h0.cmdPrint( 'ping -Q 0x30 -c 10' + h1.IP() )
info( "**** Stopping network\n" )
switch0.cmd( 'ovs-vsctl del-br dp0' )
switch0.deleteIntfs()
switch1.cmd( 'ovs-vsctl del-br dp1' )
switch1.deleteIntfs()

if __name__ == '__main__':
    setLogLevel( 'info' )
    info( "**** Scratch network demo (kernel datapath)\n" )
    Mininet.init()
    myNet()
```

21

Complex Example Given to Networking Class



- Students were asked to programmatically create the two complex topologies and interconnect them
- The left-side topology was to be deployed on one VM and right-side topology on another VM
- The Mininets were then required to communicate with each other
- This demonstrates a way to build very large and complex topologies, and test complex applications end-to-end without incurring the limitation of a single laptop's capacity
 - E.g., CAGE2 topology

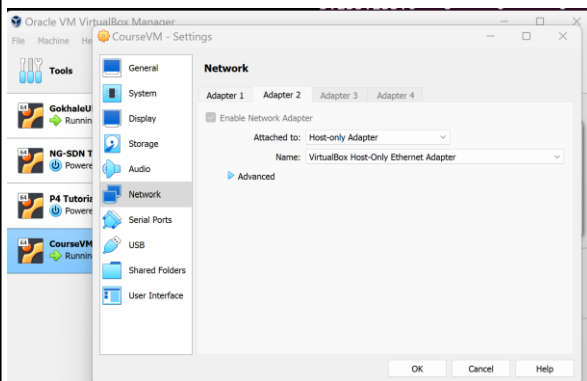
22

Approach to Making Mininets Communicate (1/7)

- Ideally, a simple solution comprising an additional routing table entry on the two VMs which host individual Mininets should work
 - This was shown to work on two VMs that used a common host-only network on the laptop
- Here is what was done to make that work
 - Start each mininet with --nat option and a -i <unique subnet> so that the hosts will get IP address from that subnet
 - Then, each VM's routing table was updated to relay traffic to the other Mininet via the underlying VM and through a specific interface
 - See screenshots next slide

23

Approach to Making Mininets Communicate (2/7)



```
gokhale@coursevm:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
```

```
gokhale@CourseVM:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
10.121.68.0 0.0.0.0 255.255.255.0 U 0 0 0 mpqemubr
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
172.18.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker_g
wbridge
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
192.168.122.0 0.0.0.0 255.255.255.0 U 0 0 0 virbr0
```

- Each VM had a second adapter for host-only network, which gets IP addresses from the 192.168.56.0/24 subnet
- This is very similar to our Chameleon VMs, which get their IP addresses in the 192.168.5.0/24 subnet
- These were the routing tables on my 2 VMs on the laptop
- Notice the rule for the 192.168.56.0 prefix, which sends packet to the network it is connected via the enp0s8 interface
 - The interface that is specified is important

24

Approach to Making Mininets Communicate (3/7)

```
gokhale@coursevm:~$ sudo mn --nat -i 10.10.1.0/24
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Warning: loopback address in /etc/resolv.conf may break host DNS over NAT
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
```

- The first VM starts a mininet with 10.10.1.0/24 subnet
- Host h1 will get 10.10.1.1, Host h2 will get 10.10.1.2 IP addresses
- We avoid 10.0.2.0/24 subnet as it is used by our NAT

```
gokhale@coursevm:~$ sudo mn --nat -i 10.10.2.0/24
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Warning: loopback address in /etc/resolv.conf may break host DNS over NAT
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
```

- The second VM starts another Mininet with 10.10.2.0/24 subnet
- Host h1 will get 10.10.2.1, Host h2 will get 10.10.2.2 IP addresses
- We avoid 10.0.2.0/24 subnet as it is used by our NAT

25

Approach to Making Mininets Communicate (4/7)

```
gokhale@coursevm:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
10.10.1.0 0.0.0.0 255.255.255.0 U 0 0 0 nat0-eth
0
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
```

- This is the routing table on VM1 after Mininet was started
- Notice the entry for 10.10.1.0, which is the address space for the hosts on that mininet

- The second VM starts another Mininet with 10.10.2.0/24 subnet
- Host h1 will get 10.10.2.1, Host h2 will get 10.10.2.2 IP addresses
- We avoid 10.0.2.0/24 subnet as it is used by our NAT

```
gokhale@coursevm:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
10.10.2.0 0.0.0.0 255.255.255.0 U 0 0 0 nat0-eth
0
10.121.68.0 0.0.0.0 255.255.255.0 U 0 0 0 mpqemubr
0
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
172.18.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker_g
wbridge
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
192.168.122.0 0.0.0.0 255.255.255.0 U 0 0 0 virbr0
```

26

Approach to Making Mininets Communicate (5/7)

```
mininet> h2 ping 10.10.2.1
PING 10.10.2.1 (10.10.2.1) 56(84) bytes of data.
```

```
mininet> h1 ping 10.10.1.2
PING 10.10.1.2 (10.10.1.2) 56(84) bytes of data.
```

- Pinging from a host of one mininet to another host of another mininet does not work but we can ping on hosts of the same mininet
- There is no route between the mininets
- The `-nat` option only allows us to go out from the boundaries of

```
enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 192.168.56.106 netmask 255.255.255.0 broadcast 192.168.56.255
inet6 fe80::f347:4108:fc28:10c6 prefixlen 64 scopeid 0x20<link>
ether 08:00:27:41:df:89 txqueuelen 1000 (Ethernet)
RX packets 1156 bytes 121153 (121.1 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 111 bytes 21002 (21.0 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 192.168.56.103 netmask 255.255.255.0 broadcast 192.168.56.255
inet6 fe80::b419:36ab:824a:ab00 prefixlen 64 scopeid 0x20<link>
ether 08:00:27:e5:28:8a txqueuelen 1000 (Ethernet)
RX packets 1559 bytes 163303 (163.3 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 129 bytes 22526 (22.5 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

- We need to add routing rules to both VMs to make this happen
- For that we need to understand our IP address on the host-only network because that is the one on which the VMs can converse
- VM1: 192.168.56.106, VM2: 192.168.156.103

27

Approach to Making Mininets Communicate (6/7)

```
gokhale@coursevm:~$ sudo ip route add 10.10.2.0/24 via 192.168.56.103 dev enp0s8
gokhale@coursevm:~$
gokhale@coursevm:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
10.10.1.0 0.0.0.0 255.255.255.0 U 0 0 0 nat0-eth
0
10.10.2.0 192.168.56.103 255.255.255.0 UG 0 0 0 enp0s8
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
```

On VM1

```
gokhale@coursevm:~$ sudo ip route add 10.10.1.0/24 via 192.168.56.106 dev enp0s8
gokhale@coursevm:~$
gokhale@coursevm:~$ route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.2.2 0.0.0.0 UG 100 0 0 enp0s3
10.0.2.0 0.0.0.0 255.255.255.0 U 100 0 0 enp0s3
10.10.1.0 192.168.56.106 255.255.255.0 UG 0 0 0 enp0s8
10.10.2.0 0.0.0.0 255.255.255.0 U 0 0 0 nat0-eth
0
10.121.68.0 0.0.0.0 255.255.255.0 U 0 0 0 mpqemubr
0
169.254.0.0 0.0.0.0 255.255.0.0 U 1000 0 0 enp0s3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker0
172.18.0.0 0.0.0.0 255.255.0.0 U 0 0 0 docker_g
wbridge
192.168.56.0 0.0.0.0 255.255.255.0 U 101 0 0 enp0s8
192.168.122.0 0.0.0.0 255.255.255.0 U 0 0 0 virbr0
```

On VM2

28

Approach to Making Mininets Communicate (7/7)

```
mininet> h2 ping 10.10.2.1
PING 10.10.2.1 (10.10.2.1) 56(84) bytes of data.
64 bytes from 10.10.2.1: icmp_seq=1 ttl=62 time=2.61 ms
64 bytes from 10.10.2.1: icmp_seq=2 ttl=62 time=3.03 ms
64 bytes from 10.10.2.1: icmp_seq=3 ttl=62 time=217 ms
64 bytes from 10.10.2.1: icmp_seq=4 ttl=62 time=3.10 ms
64 bytes from 10.10.2.1: icmp_seq=5 ttl=62 time=2.47 ms
64 bytes from 10.10.2.1: icmp_seq=6 ttl=62 time=1.76 ms
^C
--- 10.10.2.1 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5023ms
rtt min/avg/max/mdev = 1.756/38.369/217.256/80.001 ms
```

On VM1

```
mininet> h1 ping 10.10.1.2
PING 10.10.1.2 (10.10.1.2) 56(84) bytes of data.
64 bytes from 10.10.1.2: icmp_seq=1 ttl=62 time=15.7 ms
64 bytes from 10.10.1.2: icmp_seq=2 ttl=62 time=5.61 ms
64 bytes from 10.10.1.2: icmp_seq=3 ttl=62 time=2.72 ms
64 bytes from 10.10.1.2: icmp_seq=4 ttl=62 time=1.58 ms
64 bytes from 10.10.1.2: icmp_seq=5 ttl=62 time=2.29 ms
^C
--- 10.10.1.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4015ms
rtt min/avg/max/mdev = 1.582/5.577/15.690/5.238 ms
```

On VM2


29

Some Problems Encountered on Chameleon

- Although the setup is very much the same on Chameleon, there are problems getting the approach to work
- First, Chameleon has “ufw” active, which gives rise to numerous internal iptables rules
- Despite disabling ufw and flushing out all iptables rules in all tables, the solution did not work
 - Somewhere after the POSTROUTING, the reply was getting either lost or dropped
- VxLAN tunneling shows promise
 - But we could not get Open Virtual Switch-based VxLANs to work
 - So, we ended up using native Linux-based Bridging and VxLANs using the “ip” command

30

VxLAN Approach



Source: techhub.hpe.com

- VxLAN (Virtual Extensible LANs) provides an illusion of a single LAN segment to distributed resources on possibly different physical networks
- The concept is realized by tunneling Layer 2 traffic over Layer 3, i.e., ethernet traffic is carried over network traffic
- In this case, UDP is used as the protocol for tunneling
- Our responsibility is to create the tunnel endpoints and let our traffic flow through the tunnel
- This way, we can enable two or more Mininets on different VMs to communicate with each other

31

Creating VxLAN Tunnel (1/4)

- The tunnel has to be created on both VMs with each tunnel provided with an IP address (that we make up ourselves)
- The tunnel should have a unique, currently unused, ID
- In my setup, my VMs had the following IP addresses:
 - VM1: 192.168.5.115
 - VM2: 192.168.5.84
 - Your VMs will have different IP addresses; so be careful when you copy-paste these commands

32

Creating VxLAN Tunnel (2/4)

- Step 1: Create the VxLAN (this was executed on my VM2)

```
sudo ip link add vxlan0 type vxlan id 100 local 192.168.5.84 remote 192.168.5.115 dev ens3 dstport 4789
```

- where,
 - “ip link add” adds a new interface/link; must be done with “sudo”
 - vxlan0 is the name of the VxLAN and is of type “vxlan”
 - 100 is the id that we are giving this VxLAN and must be the same on the other side
 - local is the IP of the VM on which this command is executed; remote is the IP of the other VM
 - Since these IP addresses of the VM are reachable over the “ens3” interface (do ifconfig and verify), we use the “device” as “ens3)
 - VxLAN logic needs to use a UDP port, which usually is 4789
- If there are problems, you may need to add a “ufw rule”
 - sudo ufw limit 4789/udp
 - Alternately, you may just disable ufw (but take this step if things are not working)

33

Creating VxLAN Tunnel (3/4)

- Step 2: Create the VxLAN (this was executed on my VM2; issue similar command on other VM)

```
sudo ip addr add 192.168.100.2/24 dev vxlan0
```

- where,
 - Using the “ip addr add” command, we attach an arbitrary, unused, private IP address to this side of the VxLAN tunnel
 - For the other VM, use the same subnet but a different IP address
 - Effectively, we should see it show up in our ifconfig output with that ip address

34

Creating VxLAN Tunnel (4/4)

- Step 3: Activating the VxLAN (Execute this on both VMs)

```
sudo ip link set vxlan0 up
```

- where,
 - Using the “ip addr set” command, we set a property on our VxLAN
 - Here, we activate it
- Step 4: Verify creation using ifconfig or ip link show

```
sudo ip link show
```

Or simply

```
ifconfig
```

- where,
 - We should see a vxlan0 interface just like we have ens3 and others
 - There should be an IP address associated with it

35

Getting the Mininets to Communicate

- Start the Mininets on the two VMs using two separate subnets as discussed/shown before
- Use the route table manipulation steps shown earlier but this time use the IP addresses from the VxLAN and that interface
- Then test out ping, iperf and finally our application code by putting a server on one host of one of the mininets, and the client on another host of another mininet
- Similar approach can be used to deploy CASTLE workflows end-to-end

36